

SVM-hw


Ziyi Bai

2021/3/11

9.3


##(a)

```
x1 <- c(3,2,4,1,2,4,4)
x2 <- c(4,2,4,4,1,3,1)
cols <- c("red", "red", "red", "red", "blue", "blue", "blue")
plot(x1,x2,col=cols,xlim = c(0,5),ylim = c(0,5))
```




##(b)

```
plot(x1,x2,col=cols,xlim = c(0,5),ylim = c(0,5))
abline(-0.5,1)
```



##(d)

```
plot(x1,x2,col=cols,xlim = c(0,5),ylim = c(0,5))
abline(-0.5,1)
abline(-1,1,lty=2)
abline(0,1,lty=2)
```




##(e)

The support vectors for the maximal margin classifier are the points (2,1), (2,2), (4,3) and (4,4)

##(f) The 7th observations is not a support vector, so move it won't affect the outcome.

##(g)

```
plot(x1,x2,col=cols,xlim = c(0,5),ylim = c(0,5))
abline(-0.25,1)
```




9.5

##(a)

```
set.seed(9)
x1=runif(500)-0.5
x2=runif(500)-0.5
y=1*(x1^2-x2^2 > 0)
```

##(b)

```
plot(x1,x2,col=ifelse(y,"red","blue"))
```




##(c)

```
df1 <- data.frame(x1,x2,y)
fit_glm <- glm(y~x1+x2, data=df1,family = binomial)
fit_glm
```

```
##
## Call:  glm(formula = y ~ x1 + x2, family = binomial, data = df1)
##
## Coefficients:
## (Intercept)          x1          x2
##    0.05514      0.38587     -0.02653
##
## Degrees of Freedom: 499 Total (i.e. Null);  497 Residual
## Null Deviance:      692.8
## Residual Deviance: 691.2    AIC: 697.2
```

```
##(d)
```

```
pred_fit <- predict(fit_glm,data.frame(x1,x2))
plot(x1,x2,col=ifelse(pred_fit>0,"red","blue"),pch=ifelse(as.integer(pred_fit>0)==y,1,4))
```



```
##(e)
```

```
fit_glm1 <- glm(y~poly(x1,2)+poly(x2,2),data=df1,family = binomial)
```

```
## Warning: glm.fit: algorithm did not converge
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
summary(fit_glm1)
```


```
##
## Call:
## glm(formula = y ~ poly(x1, 2) + poly(x2, 2), family = binomial,
##      data = df1)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -1.079e-03 -2.000e-08 2.000e-08 2.000e-08 9.076e-04
##
## Coefficients:
##           Estimate Std. Error z value Pr(>|z|)
## (Intercept)    43.78   3063.63   0.014   0.989
## poly(x1, 2)1   1360.39 102905.10   0.013   0.989
## poly(x1, 2)2  21374.91 785951.63   0.027   0.978
## poly(x2, 2)1   -119.10   88918.85 -0.001   0.999
## poly(x2, 2)2 -21333.50 788724.67 -0.027   0.978
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 6.9276e+02 on 499 degrees of freedom
## Residual deviance: 2.4730e-06 on 495 degrees of freedom
## AIC: 10
##
## Number of Fisher Scoring iterations: 25
```

```
fit_glm2 <- glm(y~x1+x2+x1*x2,data=df1,family=binomial)
summary(fit_glm2)
```


```
##
## Call:
## glm(formula = y ~ x1 + x2 + x1 * x2, family = binomial, data = df1)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.342  -1.199   1.050   1.144   1.291
##
## Coefficients:
##           Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.05206   0.08983   0.580   0.562
## x1           0.38234   0.31036   1.232   0.218
## x2          -0.01969   0.31760  -0.062   0.951
## x1:x2        0.64537   1.12041   0.576   0.565
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 692.76 on 499 degrees of freedom
## Residual deviance: 690.87 on 496 degrees of freedom
## AIC: 698.87
##
## Number of Fisher Scoring iterations: 3
##(f)
```

```
pred_fit1 <- predict(fit_glm1, df1)
plot(x1, x2, col = ifelse(pred_fit1 > 0, "red", "blue"), pch = ifelse(as.integer(pred_fit1 > 0) == y, 1,
```




##(g)

```
df1$y <- as.factor(df1$y)
fit_svc <- svm(y~x1+x2,data=df1,kernel="linear")
pred_svc <- predict(fit_svc,df1,type="response")
plot(x1,x2,col=ifelse(pred_svc!=0,"red","blue"),pch=ifelse(pred_svc==y,1,4))
```



##(h)

```
fit_svm <- svm(y ~ x1 + x2, data = df1, kernel = "polynomial", degree = 2)
pred_svm <- predict(fit_svm, df1, type = "response")
plot(x1, x2, col = ifelse(pred_svm != 0, "red", "blue"), pch = ifelse(pred_svm == y, 1, 4))
```



##(i)

Using polynomial gives us better result.

9.7

##(a)

```
data("Auto")
Auto$Y <- ifelse(Auto$mpg > median(Auto$mpg),1,0)
Auto$Y <- as.factor(Auto$Y)
```


##(b)

```
set.seed(9)
cost <- data.frame(cost=seq(0.01,100,length.out = 10))
svm_tune <- tune(svm,Y~.,data=Auto,kernel="linear",ranges = cost)
summary(svm_tune)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
## 11.12
##
## - best performance: 0.02301282
##
## - Detailed performance results:
##   cost      error dispersion
## 1    0.01 0.07653846 0.05100638
## 2   11.12 0.02301282 0.01891104
## 3   22.23 0.03583333 0.03245677
## 4   33.34 0.03839744 0.03872235
## 5   44.45 0.03839744 0.03872235
```

```
## 6 55.56 0.03839744 0.03872235
## 7 66.67 0.03839744 0.03872235
## 8 77.78 0.03839744 0.03872235
## 9 88.89 0.03839744 0.03872235
## 10 100.00 0.03839744 0.03872235
```

```
plot(svm_tune$performances[,c(1,2)],type="l")
```



cost=11.12 has the best performance.

```
##(c)
```

```
para <- data.frame(cost=seq(0.01,100,length.out = 5),degree=seq(1,100,length.out = 5))
svm_poly <- tune(svm,Y~.,data=Auto,kernel="polynomial",ranges = para)
summary(svm_poly)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost degree
## 75.0025    1
##
## - best performance: 0.02814103
##
## - Detailed performance results:
```

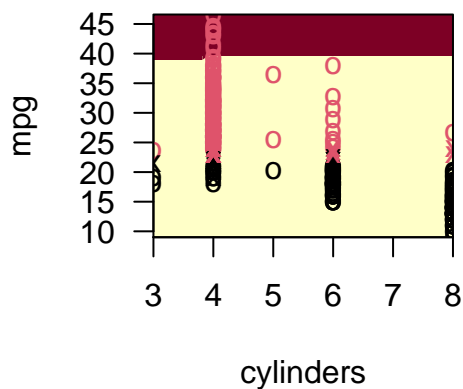
	cost	degree	error	dispersion
## 1	0.0100	1.00	0.54852564	0.01899199
## 2	25.0075	1.00	0.05102564	0.03419231
## 3	50.0050	1.00	0.03326923	0.02434857
## 4	75.0025	1.00	0.02814103	0.01893035
## 5	100.0000	1.00	0.02814103	0.01893035
## 6	0.0100	25.75	0.54852564	0.01899199


```
## 7 25.0075 25.75 0.54852564 0.01899199
## 8 50.0050 25.75 0.54852564 0.01899199
## 9 75.0025 25.75 0.54852564 0.01899199
## 10 100.0000 25.75 0.54852564 0.01899199
## 11 0.0100 50.50 0.54852564 0.01899199
## 12 25.0075 50.50 0.54852564 0.01899199
## 13 50.0050 50.50 0.54852564 0.01899199
## 14 75.0025 50.50 0.54852564 0.01899199
## 15 100.0000 50.50 0.54852564 0.01899199
## 16 0.0100 75.25 0.54852564 0.01899199
## 17 25.0075 75.25 0.54852564 0.01899199
## 18 50.0050 75.25 0.54852564 0.01899199
## 19 75.0025 75.25 0.54852564 0.01899199
## 20 100.0000 75.25 0.54852564 0.01899199
## 21 0.0100 100.00 0.54852564 0.01899199
## 22 25.0075 100.00 0.54852564 0.01899199
## 23 50.0050 100.00 0.54852564 0.01899199
## 24 75.0025 100.00 0.54852564 0.01899199
## 25 100.0000 100.00 0.54852564 0.01899199
```


##(d)

```
linear <- svm(Y ~ ., data = Auto, kernel = "linear", cost = 11.12)
polynomial <- svm(Y ~ ., data = Auto, kernel = "polynomial", cost = 100, degree = 1)
radial <- svm(Y ~ ., data = Auto, kernel = "radial", cost = 25.0075, gamma = 0.1)
pair_plot <- function(a){
  for (name in names(Auto)[!(names(Auto) %in% c("mpg", "Y", "name"))])
    plot(a, Auto, as.formula(paste("mpg~", name, sep = "")))
}
pair_plot(linear)
```



SVM classification plo





SVM classification plo





SVM classification plo





SVM classification plo





horsepower
SVM classification plo



weight
SVM classification plo




acceleration
SVM classification plo




```
pair_plot(linear)
```



SVM classification plot





SVM classification plot



SVM classification plot





SVM classification plot





```
pair_plot(polynomial)
```



SVM classification plot





SVM classification plot



SVM classification plot



SVM classification plot



##(a)

9.8

```
data("OJ")
set.seed(9)
train_oj <- sample(nrow(OJ), 800)
oj_train <- OJ[train_oj,]
oj_test <- OJ[-train_oj,]
```

```
##(b)
```

```
oj_svc <- svm(Purchase~., data=oj_train, kernel="linear", cost=0.01)
summary(oj_svc)
```

```
##
```

```
## Call:
```

```
## svm(formula = Purchase ~ ., data = oj_train, kernel = "linear", cost = 0.01)
```

```
##
```

```
##
```

```
## Parameters:
```

```
##   SVM-Type:  C-classification
```

```
##   SVM-Kernel: linear
```

```
##         cost: 0.01
```

```
##
```

```
## Number of Support Vectors: 426
```

```
##
```

```
## ( 213 213 )
```

```
##
```

```
##
```

```
## Number of Classes: 2
```

```
##
```

```
## Levels:
```

```
## CH MM
```

```
##(c)
```

```
#Training error rate
```

```
pred_train <- predict(oj_svc, oj_train)
```

```
table(pred_train, oj_train$Purchase)
```

```
##
```

```
## pred_train  CH  MM
```

```
##           CH 455  77
```

```
##           MM  50 218
```

```
#Test error rate
```

```
pred_test <- predict(oj_svc, oj_test)
```

```
table(pred_test, oj_test$Purchase)
```

```
##
```

```
## pred_test  CH  MM
```

```
##           CH 131  39
```

```
##           MM  17  83
```

```
(tr_error<- (70+61)/(428+70+61+241))
```

```
## [1] 0.16375
```

```
(te_error <- (33+21)/(143+33+21+73))
```

```
## [1] 0.2
```

```
##(d)
```

```
oj_tune <- tune(svm, Purchase~., data=oj_train, kernel="linear", ranges = data.frame(cost=seq(0.01,10, length=100)))
summary(oj_tune)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   3.34
##
## - best performance: 0.1575
##
## - Detailed performance results:
##       cost   error dispersion
## 1    0.01000 0.16625 0.03175973
## 2    0.42625 0.16125 0.03557562
## 3    0.84250 0.16125 0.03408018
## 4    1.25875 0.16125 0.03197764
## 5    1.67500 0.16125 0.03408018
## 6    2.09125 0.16000 0.03670453
## 7    2.50750 0.15875 0.03586723
## 8    2.92375 0.15875 0.03729108
## 9    3.34000 0.15750 0.03593976
## 10   3.75625 0.15750 0.03593976
## 11   4.17250 0.15750 0.03593976
## 12   4.58875 0.15750 0.03593976
## 13   5.00500 0.15875 0.03586723
## 14   5.42125 0.15875 0.03586723
## 15   5.83750 0.15875 0.03586723
## 16   6.25375 0.15875 0.03586723
## 17   6.67000 0.15875 0.03586723
## 18   7.08625 0.15875 0.03586723
## 19   7.50250 0.15875 0.03586723
## 20   7.91875 0.15875 0.03586723
## 21   8.33500 0.15875 0.03586723
## 22   8.75125 0.15875 0.03586723
## 23   9.16750 0.15875 0.03586723
## 24   9.58375 0.15875 0.03586723
## 25  10.00000 0.15875 0.03586723
```

The optimal cost is 3.75625.

##(e)

#Training error rate

```
oj_svm <- svm(Purchase ~ ., data = oj_train, kernel = "linear", cost = oj_tune$best.parameters$cost)
svm_train <- predict(oj_svm, oj_train)
table(svm_train, oj_train$Purchase)
```

```
##
## svm_train  CH  MM
##           CH 456  75
##           MM  49 220
```

```
(tr_err <- (58+64)/(425+58+64+253))
```

```
## [1] 0.1525
```



```

#Test error rate
svm_test <- predict(oj_svm, oj_test)
table(svm_test, oj_test$Purchase)

##
## svm_test  CH  MM
##          CH 131 35
##          MM  17 87

(te_err <- (27+22)/(142+27+22+79))

## [1] 0.1814815

##(f)

## radial kernel
oj_radial <- svm(Purchase~., data = oj_train, kernel="radial")
summary(oj_radial)

##
## Call:
## svm(formula = Purchase ~ ., data = oj_train, kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##         cost: 1
##
## Number of Support Vectors: 365
##
## ( 187 178 )
##
##
## Number of Classes: 2
##
## Levels:
##  CH MM

## training error rate
radial_train <- predict(oj_radial,oj_train)
table(radial_train,oj_train$Purchase)

##
## radial_train  CH  MM
##              CH 466 73
##              MM  39 222

## test error rate
radial_test <- predict(oj_radial,oj_test)
table(radial_test,oj_test$Purchase)

##
## radial_test  CH  MM
##             CH 132 37
##             MM  16 85

```

#The training error rate is 14.5% and the test error rate is 18.89%.

optimal cost

```
radial_tune <- tune(svm,Purchase~.,data=oj_train,kernel="radial",ranges = data.frame(cost=seq(0.01,10,1),
summary(radial_tune)
```

##

Parameter tuning of 'svm':

##

- sampling method: 10-fold cross validation

##

- best parameters:

cost

1.25875

##

- best performance: 0.155

##

- Detailed performance results:

cost error dispersion

1 0.01000 0.36875 0.04218428

2 0.42625 0.16125 0.02461509

3 0.84250 0.15500 0.02648375

4 1.25875 0.15500 0.03016160

5 1.67500 0.15500 0.03238227

6 2.09125 0.15625 0.03294039

7 2.50750 0.15750 0.03593976

8 2.92375 0.16125 0.03557562

9 3.34000 0.16000 0.03574602

10 3.75625 0.16125 0.03408018

11 4.17250 0.16250 0.03535534

12 4.58875 0.16625 0.03488573

13 5.00500 0.16750 0.03343734

14 5.42125 0.16625 0.03682259

15 5.83750 0.16625 0.03682259

16 6.25375 0.16625 0.03488573

17 6.67000 0.16625 0.03488573

18 7.08625 0.16625 0.03488573

19 7.50250 0.16750 0.03593976

20 7.91875 0.16875 0.03448530

21 8.33500 0.16875 0.03448530

22 8.75125 0.16875 0.03448530

23 9.16750 0.17000 0.03343734

24 9.58375 0.17000 0.03343734

25 10.00000 0.17125 0.03283481

training error rate

```
radial_svm <- svm(Purchase ~ ., data = oj_train, kernel = "radial", cost = radial_tune$best.parameters$
```

```
svm_rad <- predict(radial_svm, oj_train)
```

```
table(svm_rad, oj_train$Purchase)
```

##

svm_rad CH MM

CH 469 73

MM 36 222

```

#Test error rate
svm_rad_test <- predict(radial_svm, oj_test)
table(svm_rad_test, oj_test$Purchase)

##
## svm_rad_test  CH  MM
##           CH 132  36
##           MM  16  86

#The training error rate is 14.5% and the test error rate is 18.89%.

##(g)
## polynomial kernel
oj_poly <- svm(Purchase~.,data=oj_train,kernel="polynomial",degree=2)
summary(oj_poly)

##
## Call:
## svm(formula = Purchase ~ ., data = oj_train, kernel = "polynomial",
##      degree = 2)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##      cost:   1
##   degree:   2
##   coef.0:   0
##
## Number of Support Vectors:  431
##
## ( 219 212 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM

##(h) linear kernel with cost of 3.75625 has the best result.

```